



Applied Parallel Computing LLC  
<http://parallel-computing.pro>



CSCS

Centro Svizzero di Calcolo Scientifico  
Swiss National Supercomputing Centre

# Introduction to GPU architecture and computing

Dmitry Mikushin

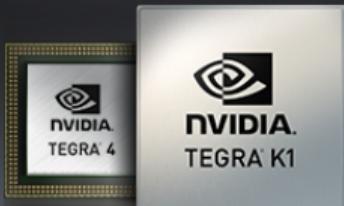
July 7, 2014

# GPU Computing is everywhere!



Courtesy of NVIDIA

# GPU Computing is everywhere!

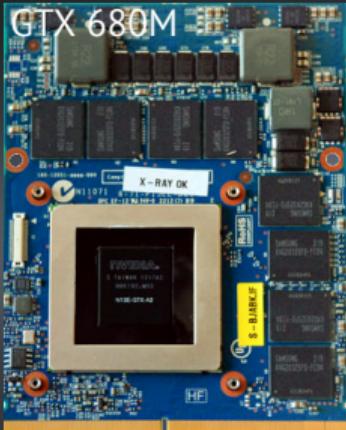


- Augmented reality
- Autonomous robots, unmanned aircrafts
- ...
- Energy-efficient data centers



Courtesy of NVIDIA

# GPU Computing is everywhere!



- Mobile GPU workstation
- GPU-accelerated CAD/CAE
- GPU apps development
- ...



Courtesy of NVIDIA

# GPU Computing is everywhere!

Tesla K40



- GPU workstation
- GPU-accelerated CAD/CAE, numerical simulations (CFD, seismic, oil&gas)
- Multiple GPUs in a single machine – MultiGPU
- ...



Courtesy of NVIDIA

# GPU Computing is everywhere!

Tesla K20X



- GPU server/cluster
- GPU-accelerated CAD/CAE, numerical simulations (CFD, seismic, oil&gas)
- Multiple GPUs in a single server – MultiGPU
- Graphics virtualization (NVIDIA GRID), remote GPU computing
- ...



Courtesy of NVIDIA

# GPU Computing is everywhere!



- Infotainment and navigation
- Advanced driver assistance systems (ADAS)
- Rear seat entertainment



Courtesy of NVIDIA

# GPU in HPC system



Oak Ridge National Laboratory Tour - Upgrading to Tesla K20 GPUs

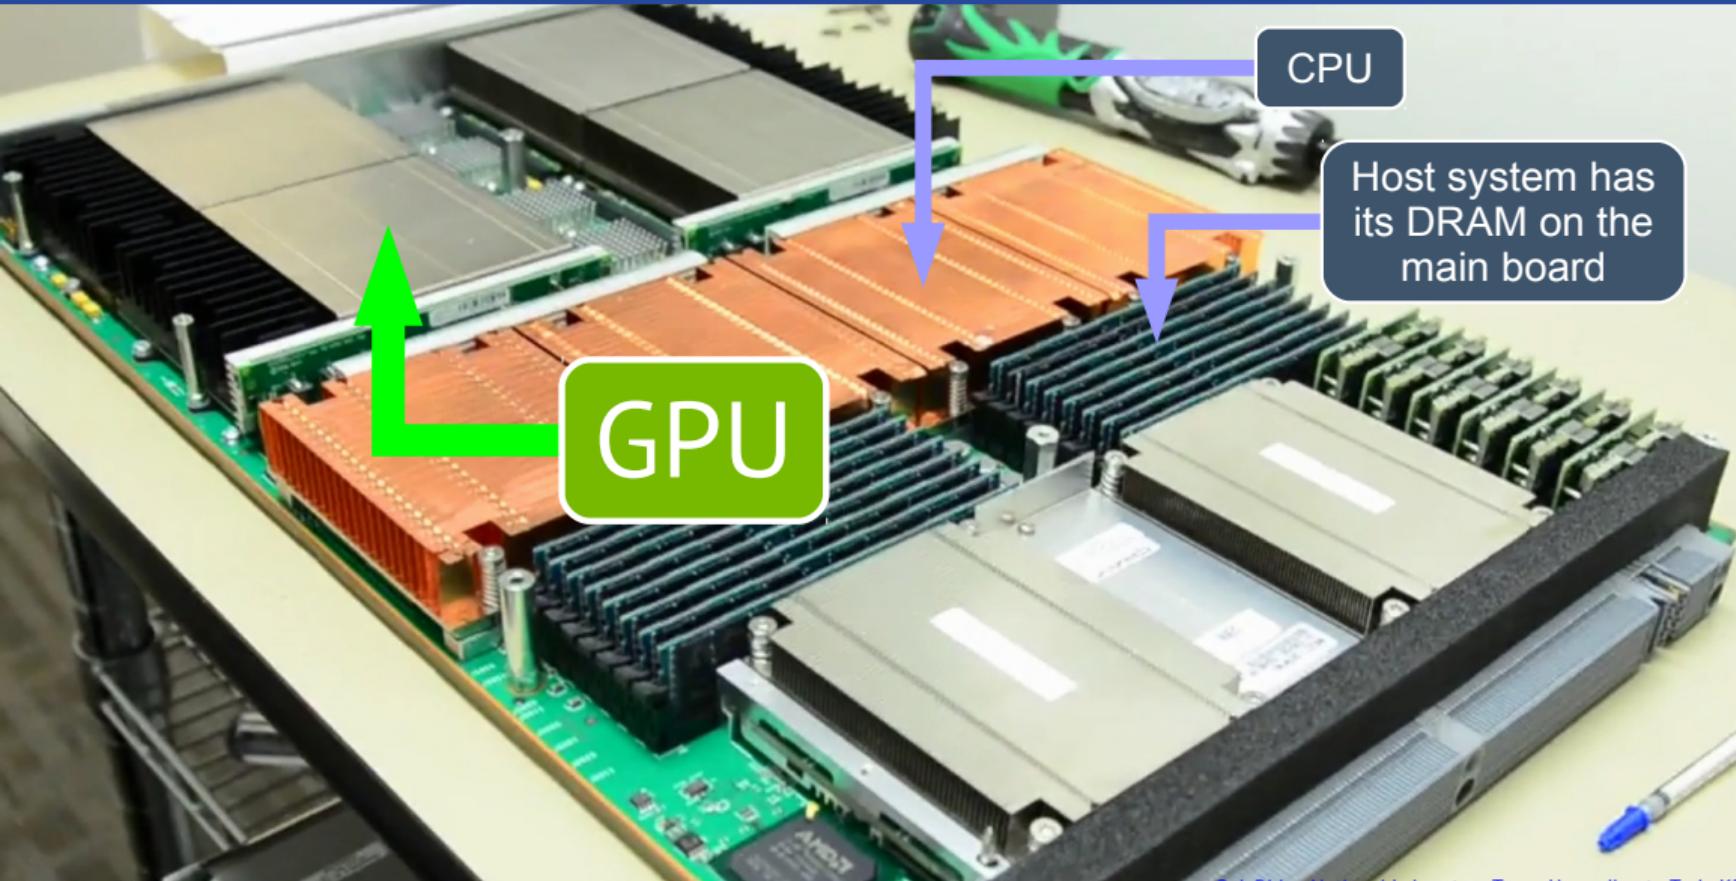
# GPU in HPC system



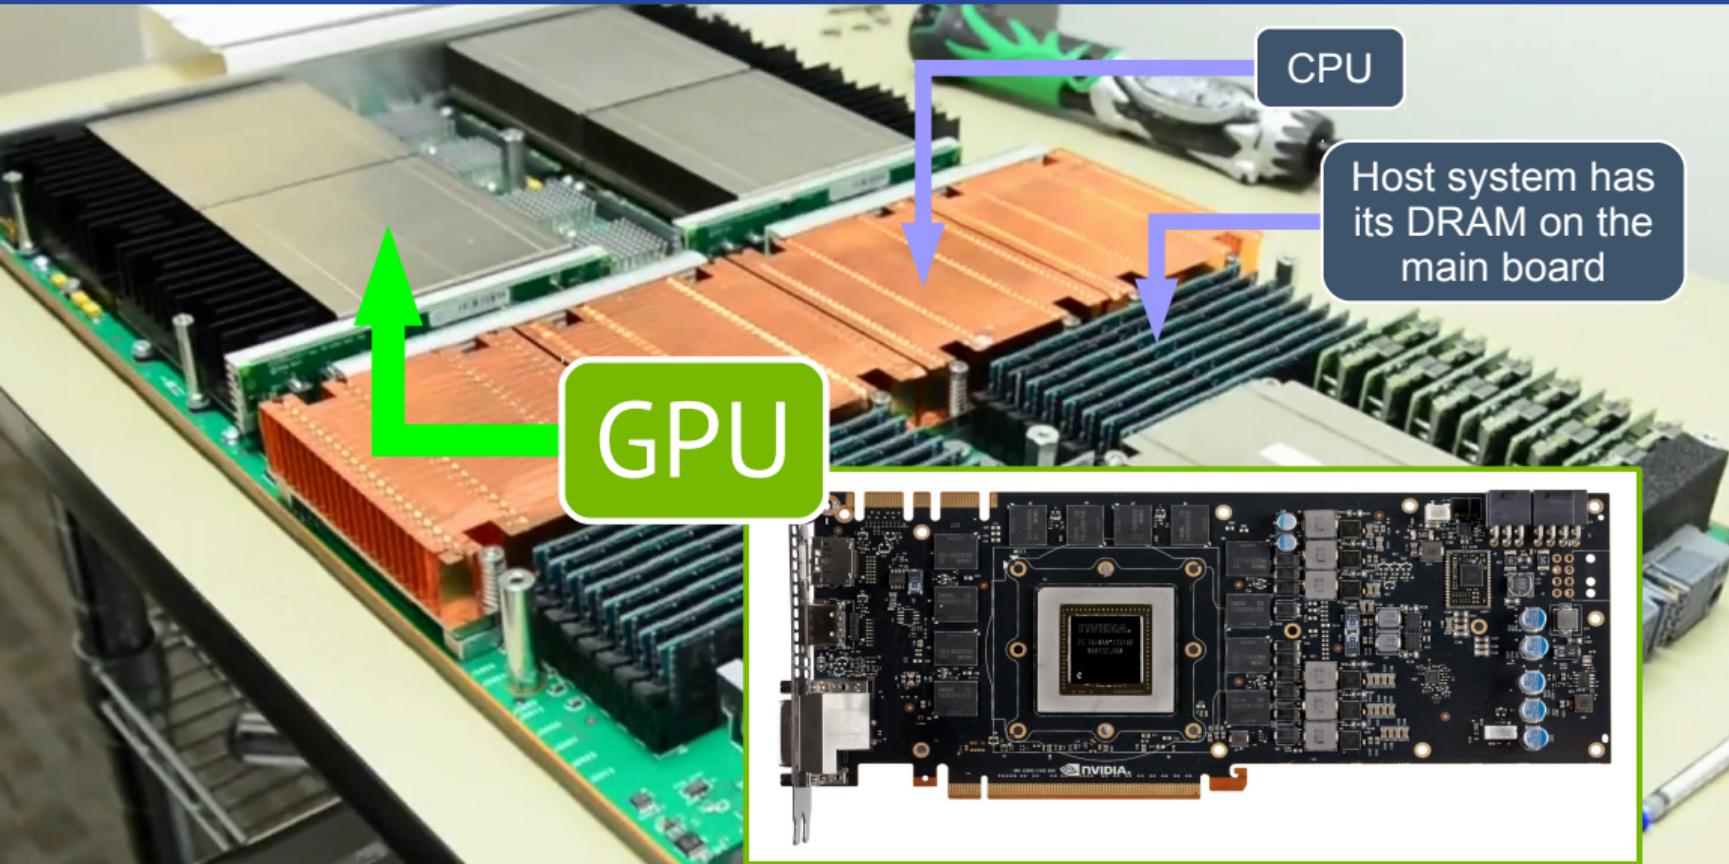
# GPU in HPC system



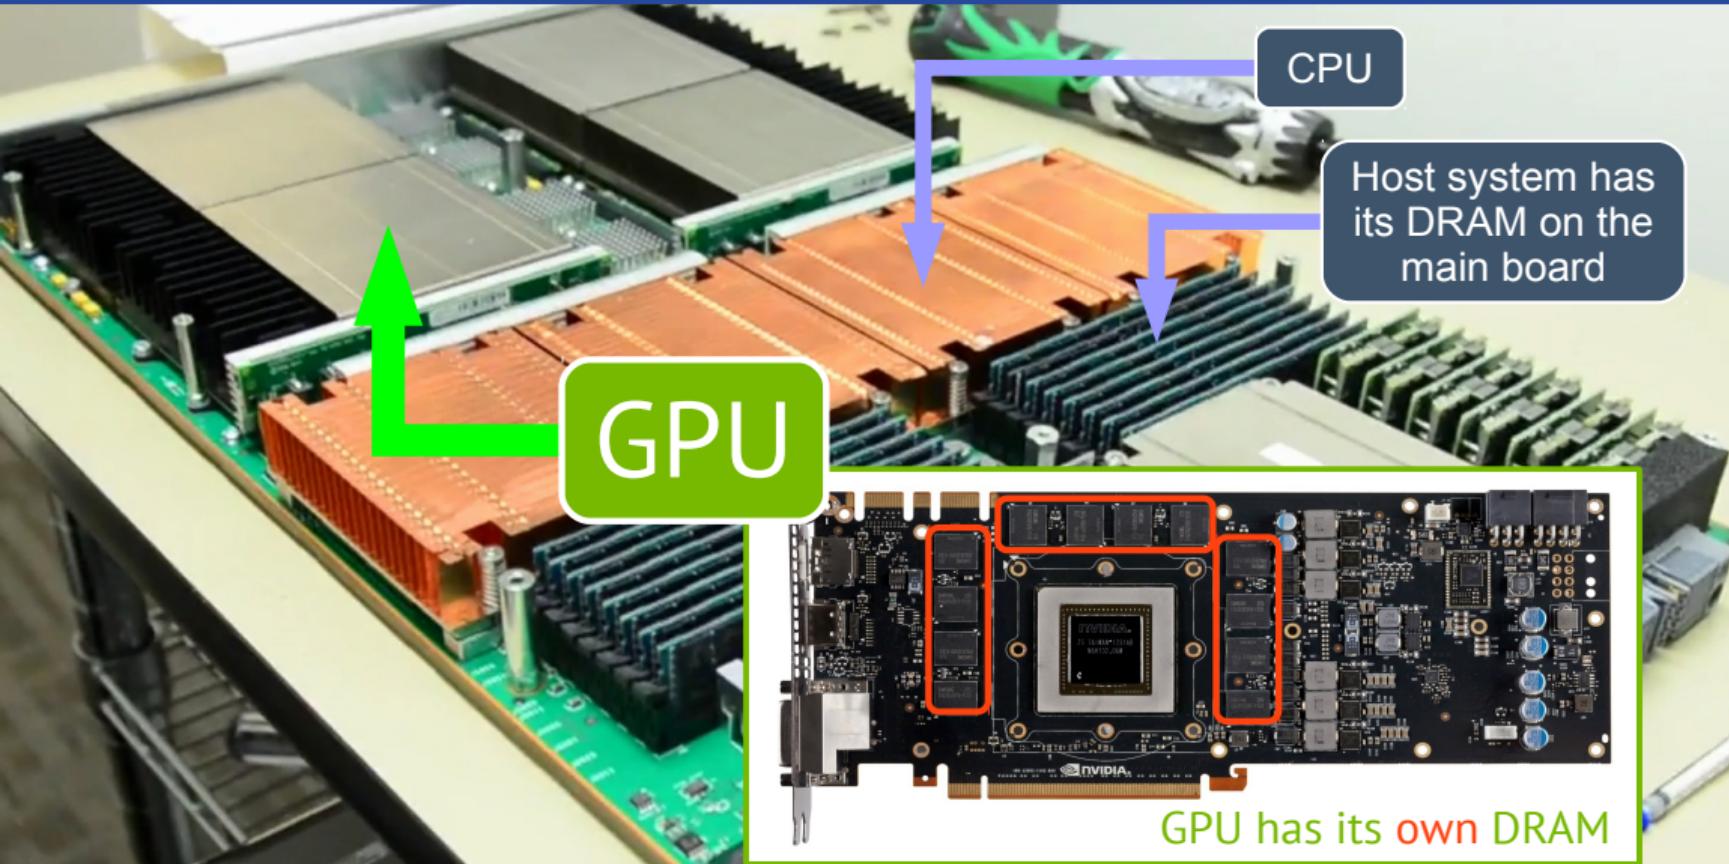
# GPU in HPC system



# GPU in HPC system

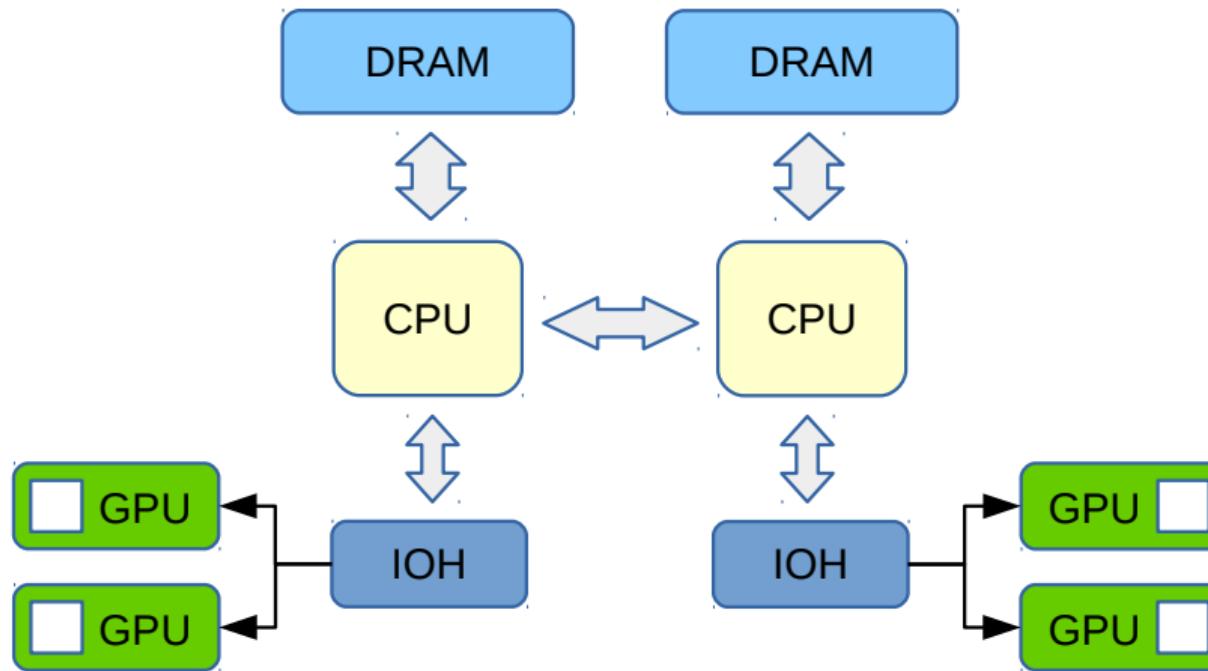


# GPU in HPC system



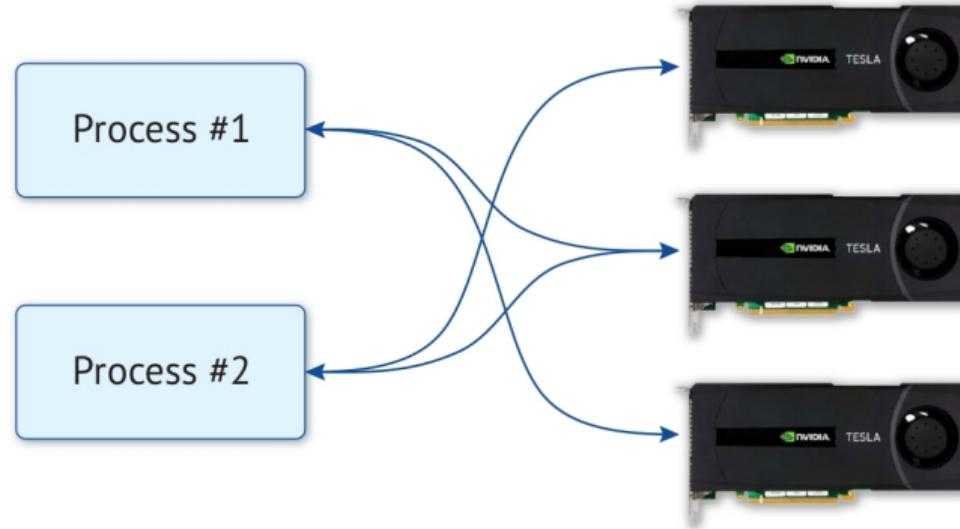
# GPU in HPC system

- MultiGPU systems are “fat” nodes: 4-8 GPUs per 2-4 CPU sockets
- Depending on I/O Hub, GPU↔CPU links might be non-uniform



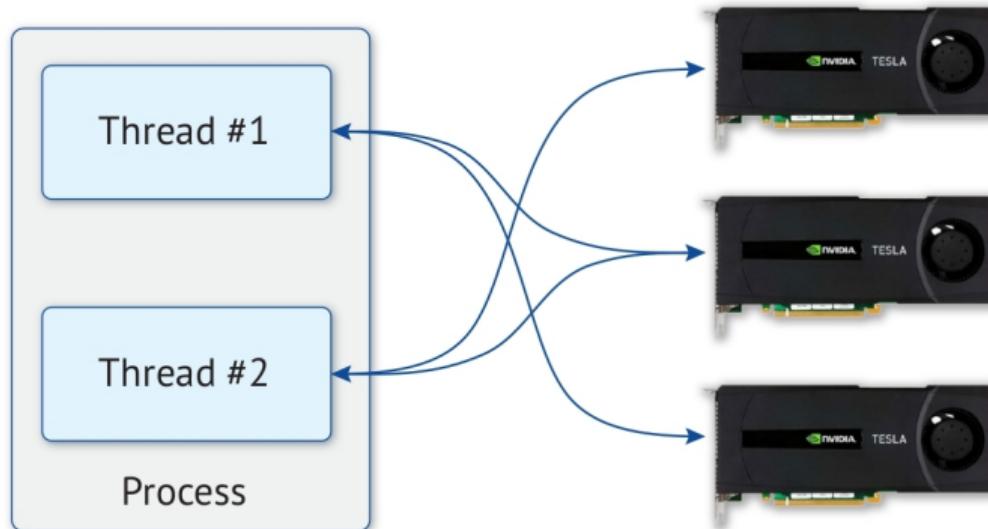
# GPU in HPC system: MPI

- Multi-process applications (e.g. MPI) can use multiple GPUs
- Single process can use multiple GPUs, multiple processes can use the same GPU (“Hyper-Q”)



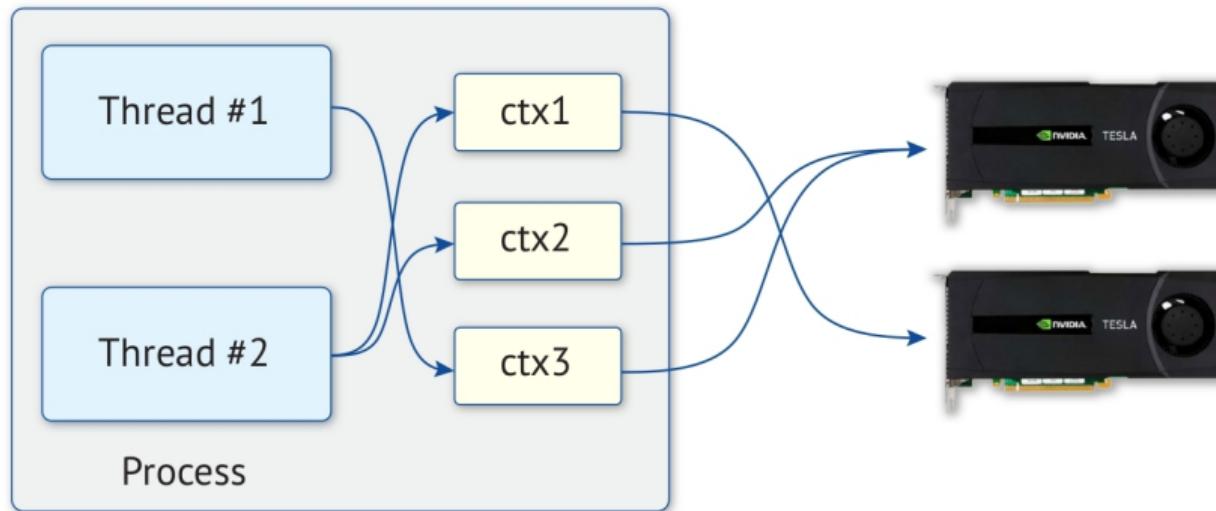
# GPU in HPC system: threads

- Multithread applications (OpenMP, POSIX, Boost, etc.) can use multiple GPUs
- Single thread can use multiple GPUs, multiple threads can use the same GPU



# GPU in HPC system: context

- Underneath any combination of threads and processes is handled by GPU *contexts*
- Thanks to contexts, GPU could be utilized in any existing MPI+OpenMP application



# Evolution of GPU ecosystem

**2008:**

- GPU computing was a technology for experts
  - Any GPU program had to be written in CUDA
  - Worked only on one specific GPU family
- 

**2014:**

- GPU computing is a technology for everyone
- A lot of free GPU-enabled libraries implement all kinds of data processing: CUBLAS, MAGMA, CUSPARSE, AmgX, CUFFT, Thrust, CUB, etc.
- Directive-based languages simplify GPU programming: OpenACC, HMPP
- Works on 5 generations of GPUs: G80 (2006), GT200 (2008), Fermi (2010), Kepler (2012), Maxwell (2014)
- CUDA is now needed only to program something very specific or outstandingly efficient

**2008:**

- GPU computing was a technology for experts
  - Any GPU program had to be written in CUDA
  - Worked only on one specific GPU family
- 

**2014:**

- GPU computing is a technology for everyone
- A lot of free GPU-enabled libraries implement all kinds of data processing: CUBLAS, MAGMA, CUSPARSE, AmgX, CUFFT, Thrust, CUB, etc.
- Directive-based languages simplify GPU programming: OpenACC, HMPP
- Works on 5 generations of GPUs: G80 (2006), GT200 (2008), Fermi (2010), Kepler (2012), Maxwell (2014)
- CUDA is now needed only to program something very specific or outstandingly efficient

# Three levels of GPU applications development



**Black box:** using GPU through a library that has GPU support underneath



**Some GPU awareness:** advice the compiler how to offload loops or data into GPU



**CUDA:** write GPU kernels manually

# Three levels of GPU applications development



**Black box:** using GPU through a library that has GPU support underneath



**Some GPU awareness:** advice the compiler how to offload loops or data into GPU



**CUDA:** write GPU kernels manually

# Three levels of GPU applications development



**Black box:** using GPU through a library that has GPU support underneath

**OpenACC®**

DIRECTIVES FOR ACCELERATORS

**Some GPU awareness:** advice the compiler how to offload loops or data into GPU



CUDA: write GPU kernels manually

# Three levels of GPU applications development



**Black box:** using GPU through a library that has GPU support underneath

**OpenACC**

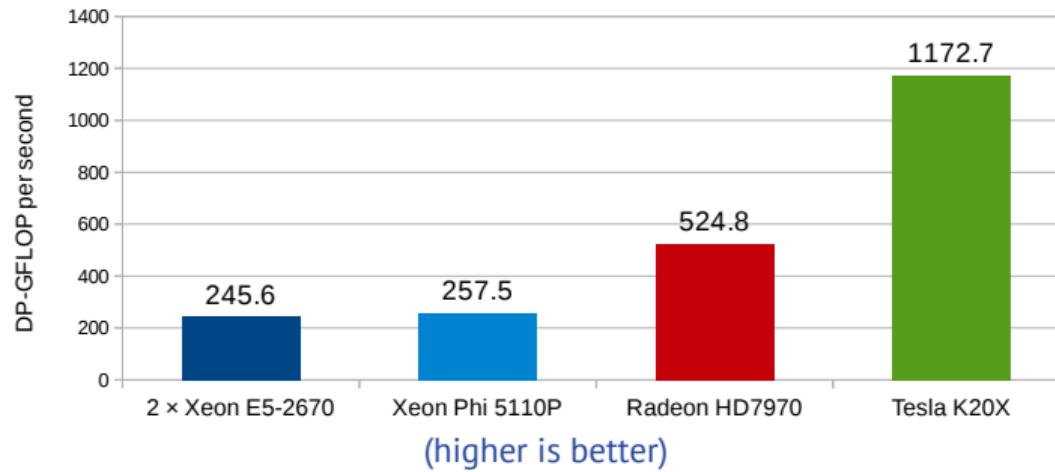
DIRECTIVES FOR ACCELERATORS

**Some GPU awareness:** advice the compiler how to offload loops or data into GPU



**CUDA:** write GPU kernels manually

# BLAS DGEMM



icpc/mkl: 2013.2.146, cublas: 5.5

# BLAS DGEMM

## CPU BLAS

```
#include <mkl.h>

double alpha = 1.0, beta = 0.0;
cblas_dgemm(CblasColMajor, CblasNoTrans, CblasNoTrans,
            n, n, n, alpha, A, n, B, n, beta, C, n);
```

## GPU-enabled BLAS

```
#include <cuda_runtime.h>
#include <cublas_v2.h>

double *A_dev = NULL, *B_dev = NULL, *C_dev = NULL;
cudaMalloc(&A_dev, n * n * sizeof(double));
cudaMalloc(&B_dev, n * n * sizeof(double));
cudaMalloc(&C_dev, n * n * sizeof(double));

cublasHandle_t handle;
cublasCreate(&handle);

cublasSetMatrix(n, n, sizeof(double), A, n, A_dev, n);
cublasSetMatrix(n, n, sizeof(double), B, n, B_dev, n);
cublasSetMatrix(n, n, sizeof(double), C, n, C_dev, n);

double alpha = 1.0, beta = 0.0;
cUBLAS_Dgemm(handle, CUBLAS_OP_N, CUBLAS_OP_N,
              n, n, n, &alpha, A_dev, n, B_dev, n, &beta, C_dev, n);
cudaDeviceSynchronize();

cUBLAS_GetMatrix(n, n, sizeof(double), C_dev, n, C, n);

cudaFree(A_dev);
cudaFree(B_dev);
cudaFree(C_dev);
cUBLAS_Destroy(handle);
```

# BLAS DGEMM

## CPU BLAS

```
#include <mkl.h>

double alpha = 1.0, beta = 0.0;
cblas_dgemm(CblasColMajor, CblasNoTrans, CblasNoTrans,
            n, n, n, alpha, A, n, B, n, beta, C, n);
```

## GPU-enabled BLAS

```
#include <cuda_runtime.h>
#include <cublas_v2.h>

double *A_dev = NULL, *B_dev = NULL, *C_dev = NULL;
cudaMalloc(&A_dev, n * n * sizeof(double));
cudaMalloc(&B_dev, n * n * sizeof(double));
cudaMalloc(&C_dev, n * n * sizeof(double));

cublasHandle_t handle;
cublasCreate(&handle);

cublasSetMatrix(n, n, sizeof(double), A, n, A_dev, n);
cublasSetMatrix(n, n, sizeof(double), B, n, B_dev, n);
cublasSetMatrix(n, n, sizeof(double), C, n, C_dev, n);

double alpha = 1.0, beta = 0.0;
cublasDgemm(handle, CUBLAS_OP_N, CUBLAS_OP_N,
            n, n, n, &alpha, A_dev, n, B_dev, n, &beta, C_dev, n);
cudaDeviceSynchronize();

cublasGetMatrix(n, n, sizeof(double), C_dev, n, C, n);

cudaFree(A_dev);
cudaFree(B_dev);
cudaFree(C_dev);
cublasDestroy(handle);
```

# BLAS DGEMM

## CPU BLAS

```
#include <mkl.h>

double alpha = 1.0, beta = 0.0;
cblas_dgemm(CblasColMajor, CblasNoTrans, CblasNoTrans,
            n, n, n, alpha, A, n, B, n, beta, C, n);
```

## GPU-enabled BLAS

```
#include <cuda_runtime.h>
#include <cUBLAS_v2.h>

double *A_dev = NULL, *B_dev = NULL, *C_dev = NULL;
cudaMalloc(&A_dev, n * n * sizeof(double));
cudaMalloc(&B_dev, n * n * sizeof(double));
cudaMalloc(&C_dev, n * n * sizeof(double));

cUBLASHandle_t handle;
cUBLASCreate(&handle);

cUBLASSetMatrix(n, n, sizeof(double), A, n, A_dev, n);
cUBLASSetMatrix(n, n, sizeof(double), B, n, B_dev, n);
cUBLASSetMatrix(n, n, sizeof(double), C, n, C_dev, n);

double alpha = 1.0, beta = 0.0;
cUBLASDgemm(handle, CUBLAS_OP_N, CUBLAS_OP_N,
             n, n, n, &alpha, A_dev, n, B_dev, n, &beta, C_dev, n);
cudaDeviceSynchronize();

cUBLASGetMatrix(n, n, sizeof(double), C_dev, n, C, n);

cudaFree(A_dev);
cudaFree(B_dev);
cudaFree(C_dev);
cUBLASDestroy(handle);
```

# BLAS DGEMM

## CPU BLAS

```
#include <mkl.h>

double alpha = 1.0, beta = 0.0;
cblas_dgemm(CblasColMajor, CblasNoTrans, CblasNoTrans,
            n, n, n, alpha, A, n, B, n, beta, C, n);
```

## GPU-enabled BLAS

```
#include <cuda_runtime.h>
#include <cublas_v2.h>

double *A_dev = NULL, *B_dev = NULL, *C_dev = NULL;
cudaMalloc(&A_dev, n * n * sizeof(double));
cudaMalloc(&B_dev, n * n * sizeof(double));
cudaMalloc(&C_dev, n * n * sizeof(double));

cublasHandle_t handle;
cublasCreate(&handle);

cublasSetMatrix(n, n, sizeof(double), A, n, A_dev, n);
cublasSetMatrix(n, n, sizeof(double), B, n, B_dev, n);
cublasSetMatrix(n, n, sizeof(double), C, n, C_dev, n);

double alpha = 1.0, beta = 0.0;
cublasDgemm(handle, CUBLAS_OP_N, CUBLAS_OP_N,
            n, n, n, &alpha, A_dev, n, B_dev, n, &beta, C_dev, n);
cudaDeviceSynchronize();

cublasGetMatrix(n, n, sizeof(double), C_dev, n, C, n);

cudaFree(A_dev);
cudaFree(B_dev);
cudaFree(C_dev);
cublasDestroy(handle);
```

# BLAS DGEMM

## CPU BLAS

```
#include <mkl.h>

double alpha = 1.0, beta = 0.0;
cblas_dgemm(CblasColMajor, CblasNoTrans, CblasNoTrans,
            n, n, n, alpha, A, n, B, n, beta, C, n);
```

## GPU-enabled BLAS

```
#include <cuda_runtime.h>
#include <cublas_v2.h>

double *A_dev = NULL, *B_dev = NULL, *C_dev = NULL;
cudaMalloc(&A_dev, n * n * sizeof(double));
cudaMalloc(&B_dev, n * n * sizeof(double));
cudaMalloc(&C_dev, n * n * sizeof(double));

cublasHandle_t handle;
cublasCreate(&handle);

cublasSetMatrix(n, n, sizeof(double), A, n, A_dev, n);
cublasSetMatrix(n, n, sizeof(double), B, n, B_dev, n);
cublasSetMatrix(n, n, sizeof(double), C, n, C_dev, n);

double alpha = 1.0, beta = 0.0;
cUBLAS_Dgemm(handle, CUBLAS_OP_N, CUBLAS_OP_N,
              n, n, n, &alpha, A_dev, n, B_dev, n, &beta, C_dev, n);
cudaDeviceSynchronize();

cUBLAS_GetMatrix(n, n, sizeof(double), C_dev, n, C, n);

cudaFree(A_dev);
cudaFree(B_dev);
cudaFree(C_dev);
cUBLAS_Destroy(handle);
```

# BLAS DGEMM

## CPU BLAS

```
#include <mkl.h>

double alpha = 1.0, beta = 0.0;
cblas_dgemm(CblasColMajor, CblasNoTrans, CblasNoTrans,
            n, n, n, alpha, A, n, B, n, beta, C, n);
```

## GPU-enabled BLAS

```
#include <cuda_runtime.h>
#include <cublas_v2.h>

double *A_dev = NULL, *B_dev = NULL, *C_dev = NULL;
cudaMalloc(&A_dev, n * n * sizeof(double));
cudaMalloc(&B_dev, n * n * sizeof(double));
cudaMalloc(&C_dev, n * n * sizeof(double));

cublasHandle_t handle;
cublasCreate(&handle);

cublasSetMatrix(n, n, sizeof(double), A, n, A_dev, n);
cublasSetMatrix(n, n, sizeof(double), B, n, B_dev, n);
cublasSetMatrix(n, n, sizeof(double), C, n, C_dev, n);

double alpha = 1.0, beta = 0.0;
cUBLAS_Dgemm(handle, CUBLAS_OP_N, CUBLAS_OP_N,
              n, n, n, &alpha, A_dev, n, B_dev, n, &beta, C_dev, n);
cudaDeviceSynchronize();

cUBLAS_GetMatrix(n, n, sizeof(double), C_dev, n, C, n);

cudaFree(A_dev);
cudaFree(B_dev);
cudaFree(C_dev);
cUBLAS_Destroy(handle);
```

# BLAS DGEMM

## CPU BLAS

```
#include <mkl.h>

double alpha = 1.0, beta = 0.0;
cblas_dgemm(CblasColMajor, CblasNoTrans, CblasNoTrans,
            n, n, n, alpha, A, n, B, n, beta, C, n);
```

## GPU-enabled BLAS

```
#include <cuda_runtime.h>
#include <cublas_v2.h>

double *A_dev = NULL, *B_dev = NULL, *C_dev = NULL;
cudaMalloc(&A_dev, n * n * sizeof(double));
cudaMalloc(&B_dev, n * n * sizeof(double));
cudaMalloc(&C_dev, n * n * sizeof(double));

cublasHandle_t handle;
cublasCreate(&handle);

cublasSetMatrix(n, n, sizeof(double), A, n, A_dev, n);
cublasSetMatrix(n, n, sizeof(double), B, n, B_dev, n);
cublasSetMatrix(n, n, sizeof(double), C, n, C_dev, n);

double alpha = 1.0, beta = 0.0;
cUBLAS_Dgemm(handle, CUBLAS_OP_N, CUBLAS_OP_N,
              n, n, n, &alpha, A_dev, n, B_dev, n, &beta, C_dev, n);
cudaDeviceSynchronize();

cUBLAS_GetMatrix(n, n, sizeof(double), C_dev, n, C, n);

cudaFree(A_dev);
cudaFree(B_dev);
cudaFree(C_dev);
cUBLAS_Destroy(handle);
```

## CPU BLAS

```
#include <mkl.h>
```

```
double alpha = 1.0, beta = 0.0;  
cblas_dgemm(CblasColMajor, CblasNoTrans, CblasNoTrans,  
            n, n, n, alpha, A, n, B, n, beta, C, n);
```

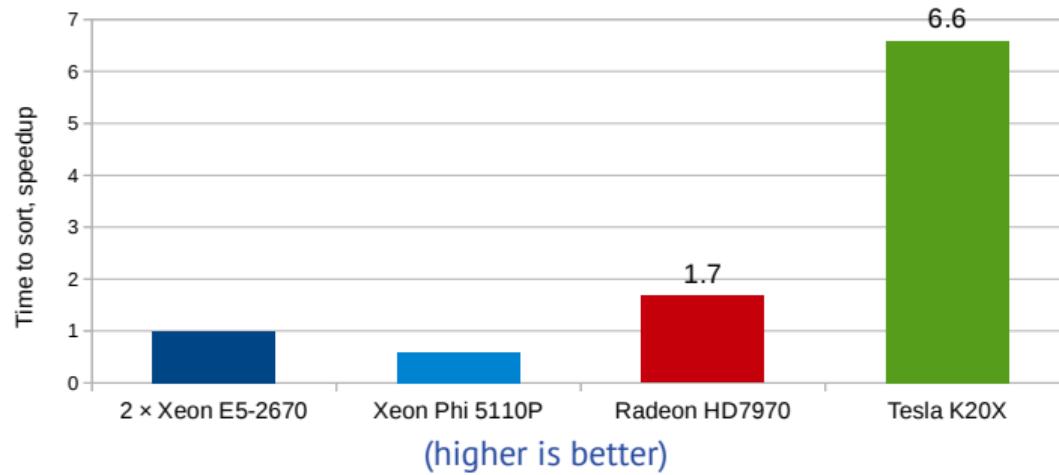
## GPU-enabled BLAS

CUDA 6.0 introduces **NVBLAS** library:

- Automatically routes standard BLAS3 calls to CUBLAS (by preloading libnvblas.so)
- Can spread work across multiple GPUs (cuBLAS-XT)
- Usable with any app that calls BLAS3 functions (Octave, Scilab, etc.)

```
double alpha = 1.0, beta = 0.0;  
dgemm_('n', 'n',  
        &n, &n, &n, &alpha, A, &n, B, &n, &beta, C, &n);
```

# Sorting 128M key-value pairs



icpc: 2013.2.146, tbb: 4.1, thrust: 5.5

# Sorting 128M key-value pairs

```
#include <thrust/device_vector.h>
#include <thrust/generate.h>
#include <thrust/sort.h>
#include <thrust/copy.h>

// Generate the random keys and values on the host.
host_vector<float> h_keys(n);
host_vector<float> h_vals(n);
for (int i = 0; i < n; i++)
{
    h_keys[i] = drand48();
    h_vals[i] = drand48();
}

// Transfer data to the device.
device_vector<float> d_keys = h_keys;
device_vector<float> d_vals = h_vals;

// Sort!
sort_by_key(d_keys.begin(), d_keys.end(), d_vals.begin());
cudaDeviceSynchronize();

// Transfer data back to host.
copy(d_keys.begin(), d_keys.end(), h_keys.begin());
copy(d_vals.begin(), d_vals.end(), h_vals.begin());
```

# Sorting 128M key-value pairs

```
#include <thrust/device_vector.h>
#include <thrust/generate.h>
#include <thrust/sort.h>
#include <thrust/copy.h>

// Generate the random keys and values on the host.
host_vector<float> h_keys(n);
host_vector<float> h_vals(n);
for (int i = 0; i < n; i++)
{
    h_keys[i] = drand48();
    h_vals[i] = drand48();
}

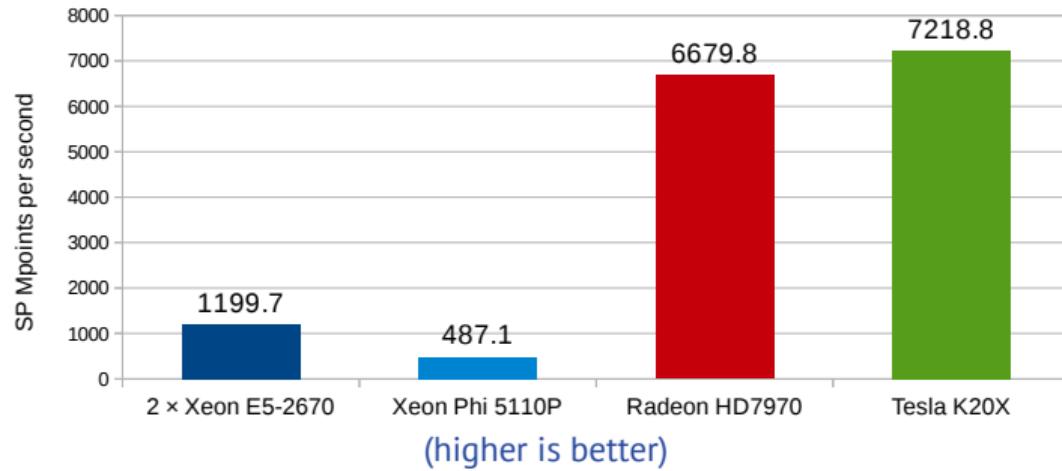
// Transfer data to the device.
device_vector<float> d_keys = h_keys;
device_vector<float> d_vals = h_vals;

// Sort!
sort_by_key(d_keys.begin(), d_keys.end(), d_vals.begin());
cudaDeviceSynchronize();

// Transfer data back to host.
copy(d_keys.begin(), d_keys.end(), h_keys.begin());
copy(d_vals.begin(), d_vals.end(), h_vals.begin());
```

Relatively new [CUB](#) library from NVResearch offers radix sort with **2.5×** higher perf than Thrust!

# Wave propagation stencil



icpc: 2013.2.146, cublas: 5.5

# Wave propagation stencil

```
void wave13pt(const int nx, const int ny, const int ns,
    const real m0, const real m1, const real m2,
    const real* const __restrict__ w0, const real* const __restrict__ w1,
    real* const __restrict__ w2)
{
    size_t szarray = (size_t)nx * ny * ns;
    #pragma acc kernels loop independent gang(ns), present(w0[0:szarray], w1[0:szarray], w2[0:szarray])
    for (int k = 2; k < ns - 2; k++)
    {
        #pragma acc loop independent
        for (int j = 2; j < ny - 2; j++)
        {
            #pragma acc loop independent vector(512)
            for (int i = 2; i < nx - 2; i++)
            {
                _A(w2, k, j, i) = m0 * _A(w1, k, j, i) - _A(w0, k, j, i) +
                    m1 * (
                        _A(w1, k, j, i+1) + _A(w1, k, j, i-1) +
                        _A(w1, k, j+1, i) + _A(w1, k, j-1, i) +
                        _A(w1, k+1, j, i) + _A(w1, k-1, j, i)) +
                    m2 * (
                        _A(w1, k, j, i+2) + _A(w1, k, j, i-2) +
                        _A(w1, k, j+2, i) + _A(w1, k, j-2, i) +
                        _A(w1, k+2, j, i) + _A(w1, k-2, j, i));
            }
        }
    }
}
```

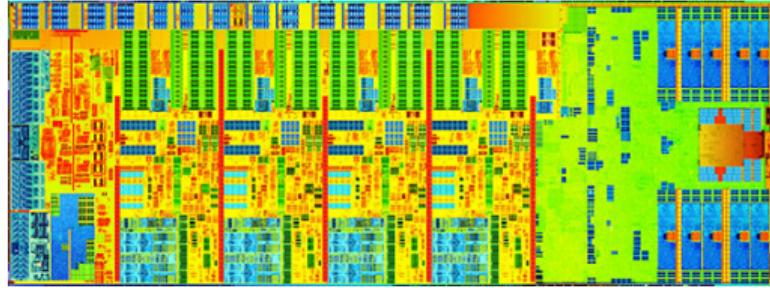
# Wave propagation stencil

```
void wave13pt(const int nx, const int ny, const int ns,
    const real m0, const real m1, const real m2,
    const real* const __restrict__ w0, const real* const __restrict__ w1,
    real* const __restrict__ w2)
{
    size_t szarray = (size_t)nx * ny * ns;
#pragma acc kernels loop independent gang(ns), present(w0[0:szarray], w1[0:szarray], w2[0:szarray])
    for (int k = 2; k < ns - 2; k++)
    {
        #pragma acc loop independent
        for (int j = 2; j < ny - 2; j++)
        {
            #pragma acc loop independent vector(512)
            for (int i = 2; i < nx - 2; i++)
            {
                _A(w2, k, j, i) = m0 * _A(w1, k, j, i) - _A(w0, k, j, i) +
                    m1 * (
                        _A(w1, k, j, i+1) + _A(w1, k, j, i-1) +
                        _A(w1, k, j+1, i) + _A(w1, k, j-1, i) +
                        _A(w1, k+1, j, i) + _A(w1, k-1, j, i)) +
                    m2 * (
                        _A(w1, k, j, i+2) + _A(w1, k, j, i-2) +
                        _A(w1, k, j+2, i) + _A(w1, k, j-2, i) +
                        _A(w1, k+2, j, i) + _A(w1, k-2, j, i));
            }
        }
    }
}
```

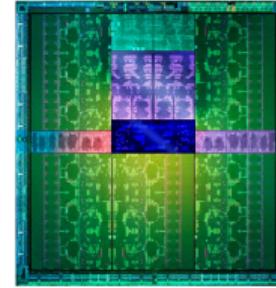
OpenACC: offload stencil computations to GPU with high efficiency:

- PGI OpenACC 14.1 often generates faster kernels than hand-written CUDA

# CPU architecture **vs** GPU architecture

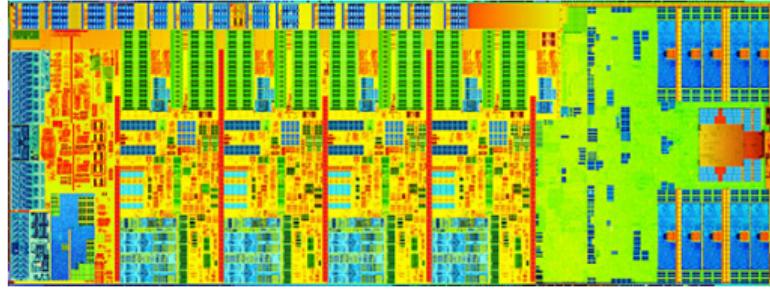


VS

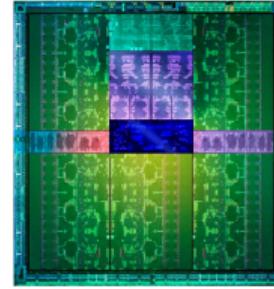


- Several complex cores
- SIMD vectorization
- Large caches
- Thousands of simple compute units
- Thousands of scalar cores
- Large register file + shared memory

# CPU architecture **vs** GPU architecture

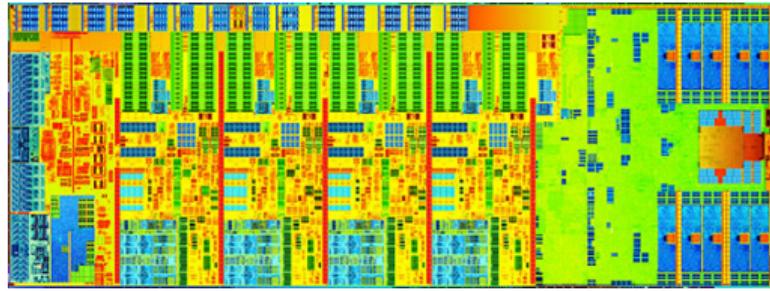


VS

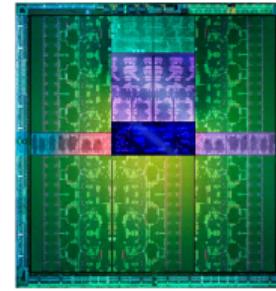


- Several complex cores
- SIMD vectorization
- Large caches
- Thousands of simple compute units
- Thousands of scalar cores
- Large register file + shared memory

# CPU architecture **vs** GPU architecture

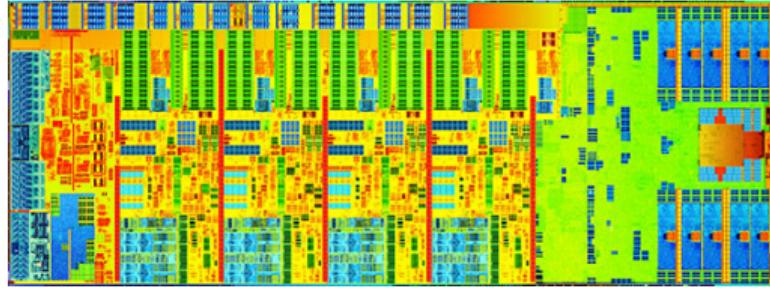


VS

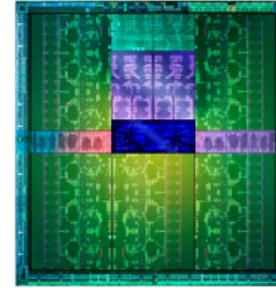


- Several complex cores
  - SIMD vectorization
  - Large caches
- 
- Thousands of simple compute units
  - Thousands of scalar cores
  - Large register file + shared memory

# CPU architecture **vs** GPU architecture

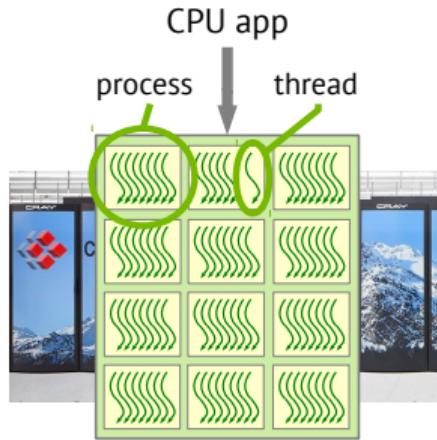


VS

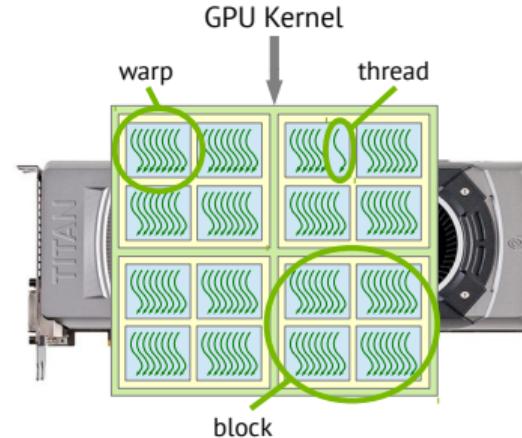


- Several complex cores
- SIMD vectorization
- Large caches
- Thousands of simple compute units
- Thousands of scalar cores
- Large register file + shared memory

# CPU execution model **vs** GPU execution model



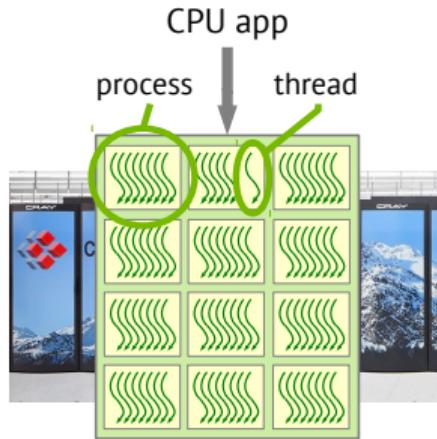
VS



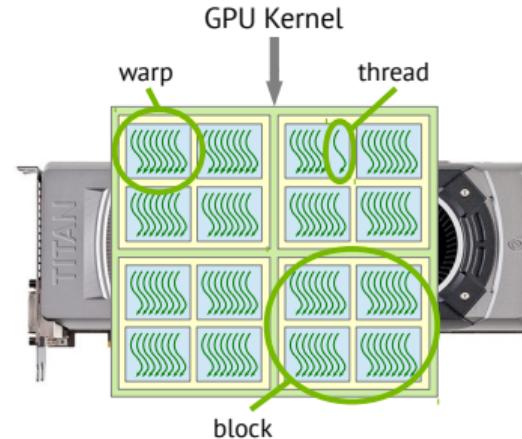
- Processes & threads maintained by OS
- Memory latency is relaxed with 3 levels of caches

- *Lightweight threads* maintained by GPU driver
- Memory latency is relaxed by active threads switching

# CPU execution model **vs** GPU execution model



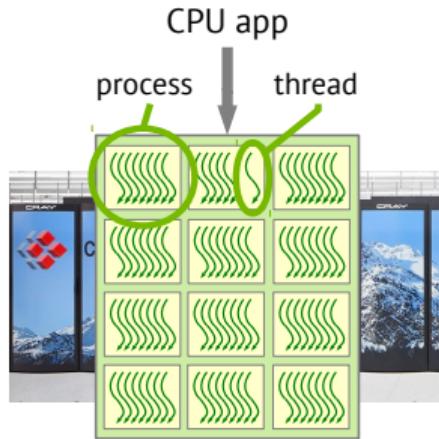
VS



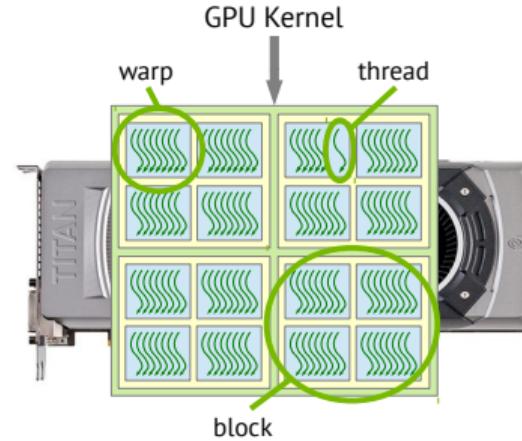
- Processes & threads maintained by OS
- Memory latency is relaxed with 3 levels of caches

- *Lightweight threads* maintained by GPU driver
- Memory latency is relaxed by active threads switching

# CPU execution model **vs** GPU execution model



VS



- Processes & threads maintained by OS
- Memory latency is relaxed with 3 levels of caches

- *Lightweight* threads maintained by GPU driver
- Memory latency is relaxed by active threads switching

# What makes the GPU application efficient?



GPU application

GPGPU computations are scalar  $\Rightarrow$  no hard time with vectorization!

- 1 The GPU task should be partitioned into significant amount of small parallel tasks, to be processed by GPU threads
- 2 Task threads should maintain the *balanced use* of two key resources: registers and shared memory
- 3 Task threads altogether should access consecutive memory addresses (coalescing)
- 4 Task threads should not do divergent if...else branching

# What makes the GPU application efficient?



GPU application

GPGPU computations are scalar  $\Rightarrow$  no hard time with vectorization!

- 1 The GPU task should be partitioned into significant amount of small parallel tasks, to be processed by GPU threads
- 2 Task threads should maintain the *balanced use* of two key resources: registers and shared memory
- 3 Task threads altogether should access consecutive memory addresses (coalescing)
- 4 Task threads should not do divergent if...else branching

# What makes the GPU application efficient?



Partitioned into small parallel tasks

GPGPU computations are scalar  $\Rightarrow$  no hard time with vectorization!

- 1 The GPU task should be partitioned into significant amount of small parallel tasks, to be processed by GPU threads
- 2 Task threads should maintain the *balanced use* of two key resources: registers and shared memory
- 3 Task threads altogether should access consecutive memory addresses (coalescing)
- 4 Task threads should not do divergent if...else branching

# What makes the GPU application efficient?



Partitioned into small parallel tasks

GPGPU computations are scalar  $\Rightarrow$  no hard time with vectorization!

- 1 The GPU task should be partitioned into significant amount of small parallel tasks, to be processed by GPU threads
- 2 Task threads should maintain the *balanced use* of two key resources: registers and shared memory
- 3 Task threads altogether should access consecutive memory addresses (coalescing)
- 4 Task threads should not do divergent if...else branching

# What makes the GPU application efficient?



Partitioned into small parallel tasks

GPGPU computations are scalar  $\Rightarrow$  no hard time with vectorization!

- 1 The GPU task should be partitioned into significant amount of small parallel tasks, to be processed by GPU threads
- 2 Task threads should maintain the *balanced use* of two key resources: registers and shared memory
- 3 Task threads altogether should access consecutive memory addresses (coalescing)
- 4 Task threads should not do divergent if...else branching

# What makes the GPU application efficient?

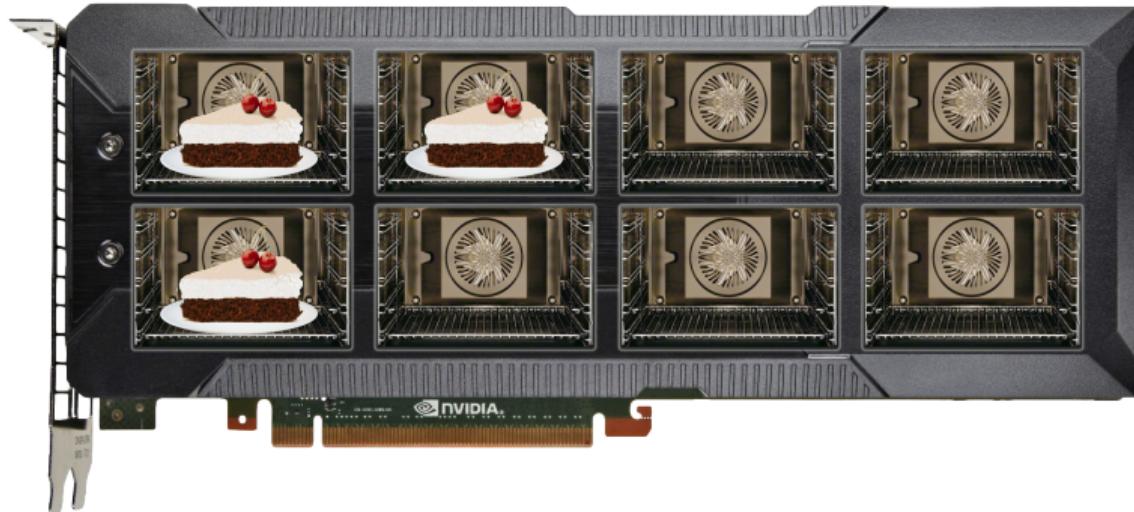


Partitioned into small parallel tasks

GPGPU computations are scalar  $\Rightarrow$  no hard time with vectorization!

- 1 The GPU task should be partitioned into significant amount of small parallel tasks, to be processed by GPU threads
- 2 Task threads should maintain the *balanced use* of two key resources: registers and shared memory
- 3 Task threads altogether should access consecutive memory addresses (coalescing)
- 4 Task threads should not do divergent if...else branching

- 1 What if the amount of parallel tasks is too small?



⇒ Compute resources are wasted

# GPU efficiency overview

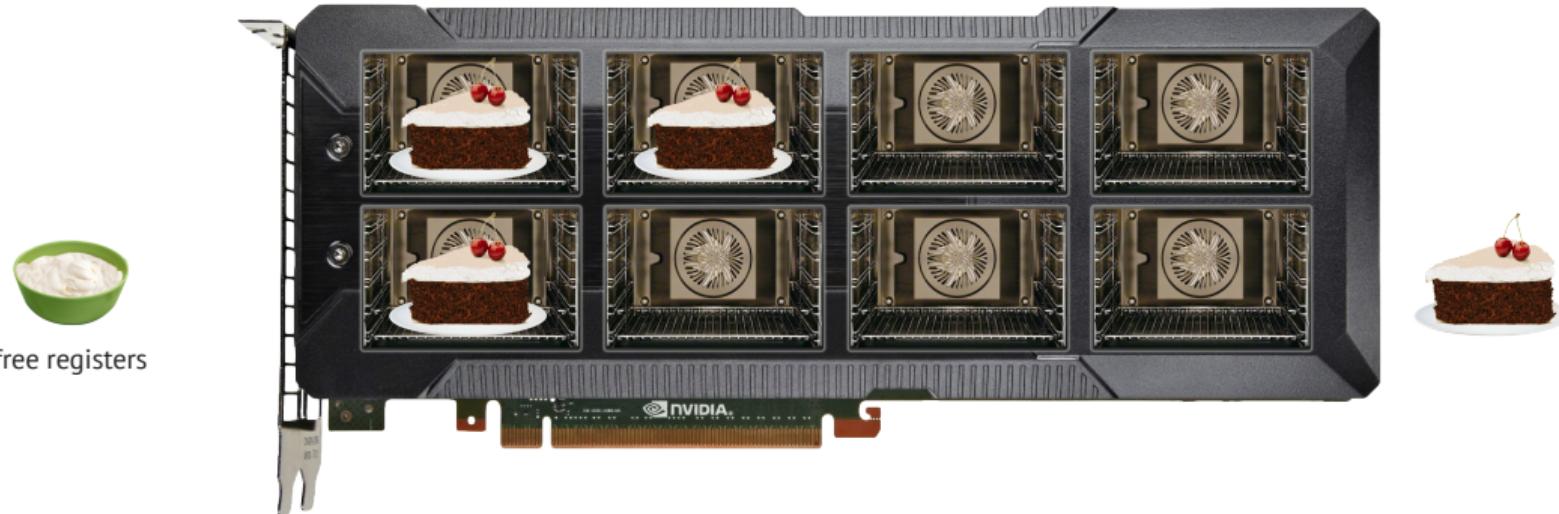
- 1 It's fine if amount of tasks is larger than the number of compute cores:



⇒ In this case the remaining tasks will be automatically processed after finishing with the first group

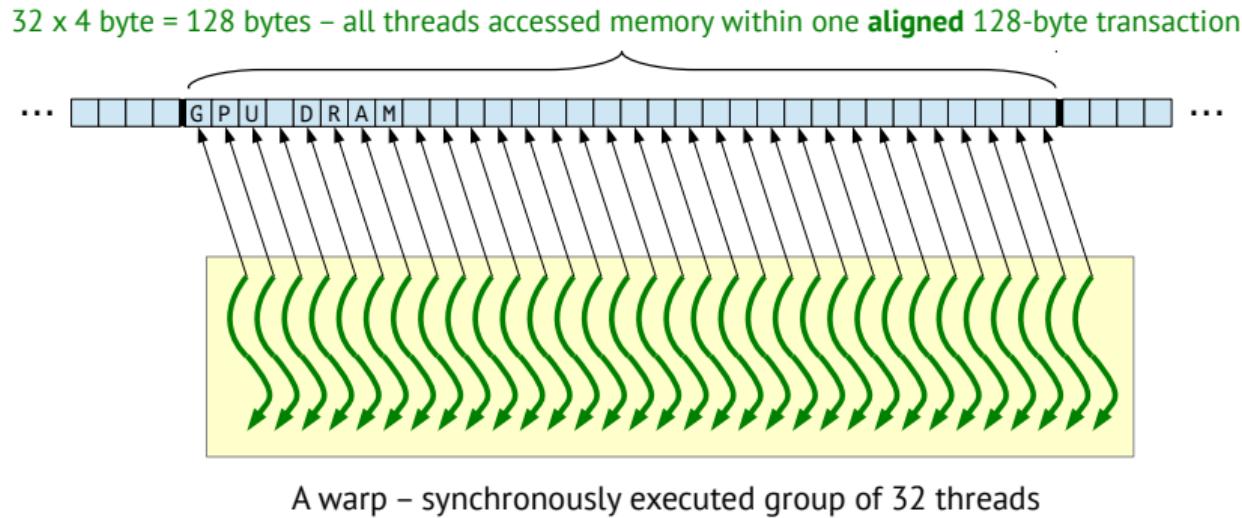
# GPU efficiency overview

- 2 If the use of registers or shared memory is imbalanced:



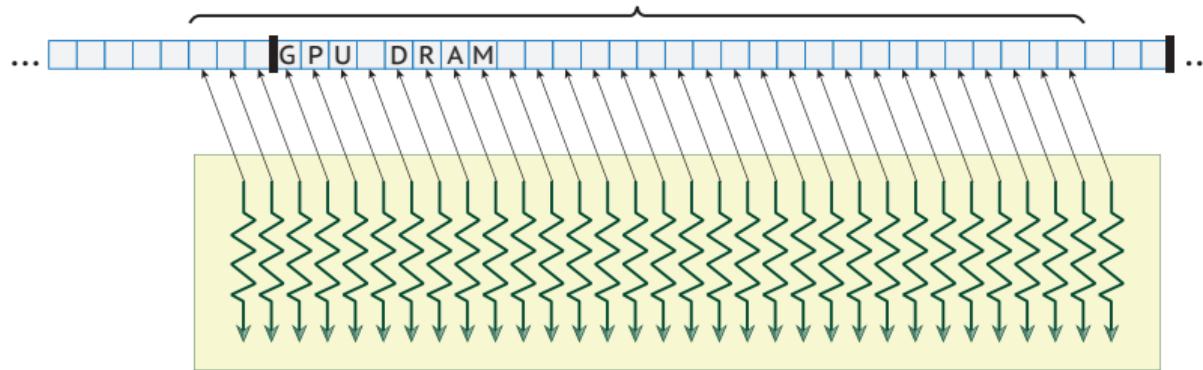
Although the free compute and shared memory resources exist, smaller amount of tasks already used all available shared memory  $\Rightarrow$  other tasks have to wait, compute resources are wasted

- 3 Task threads altogether should access consecutive memory addresses (coalescing):



- 3 Task threads altogether should access consecutive memory addresses (coalescing):

2x 128-byte transactions due to misalignment: 12 bytes from first, 116 bytes from second

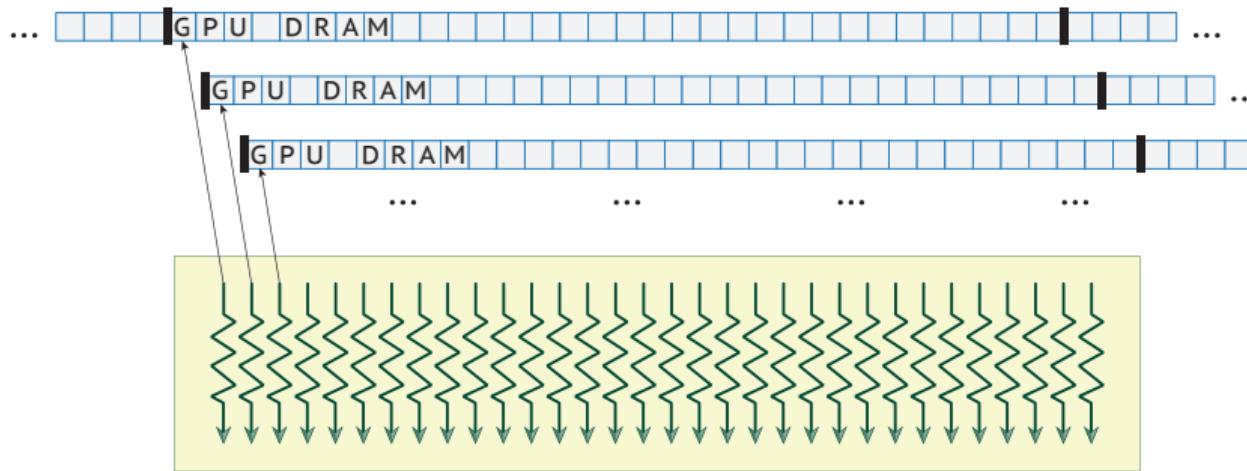


A warp – synchronously executed group of 32 threads

# GPU efficiency overview

- 3 Task threads altogether should access consecutive memory addresses (coalescing):

disaster: each thread loads 128-byte transaction to use only 4 bytes of it  
(very low memory efficiency!)



A warp – synchronously executed group of 32 threads

## 4 Task threads should not do divergent if...else branching

Threads, whose flag is false will do nothing  
(stall)  $\Rightarrow$  no divergence, just some waste of  
compute resource

```
if (flags[threadIdx.x])
{
    // A lot of code
}
```

Each thread will do its branch and then  
stall while some other thread is doing alter-  
native branch (if there is at least one  
such thread)  $\Rightarrow$  divergent branching

```
if (flags[threadIdx.x])
{
    // A lot of code
}
else
{
    // A lot of code
}
```

# Conclusion

- 1 GPUs offer much higher performance of massively-parallel computations
- 2 You can start using GPUs in your application very easily:
  - Seamless integration into applications with existing parallelism (MPI, OpenMP)
  - Fairly efficient GPU-enabled libraries exist for many types of computational problems
  - Nowadays CUDA is needed only to develop something very specific and/or highly-tuned
- 3 GPU architecture and execution model are designed for massive parallelism
- 4 GPU efficiency is based on 4 factors:
  - Sufficient parallelism
  - Balanced use of registers and shared memory
  - Efficient DRAM utilization
  - Minimization of divergent branching

# Conclusion

- 1 GPUs offer much higher performance of massively-parallel computations
- 2 You can start using GPUs in your application very easily:
  - Seamless integration into applications with existing parallelism (MPI, OpenMP)
  - Fairly efficient GPU-enabled libraries exist for many types of computational problems
  - Nowadays CUDA is needed only to develop something very specific and/or highly-tuned
- 3 GPU architecture and execution model are designed for massive parallelism
- 4 GPU efficiency is based on 4 factors:
  - Sufficient parallelism
  - Balanced use of registers and shared memory
  - Efficient DRAM utilization
  - Minimization of divergent branching

# Conclusion

- 1 GPUs offer much higher performance of massively-parallel computations
- 2 You can start using GPUs in your application very easily:
  - Seamless integration into applications with existing parallelism (MPI, OpenMP)
  - Fairly efficient GPU-enabled libraries exist for many types of computational problems
  - Nowadays CUDA is needed only to develop something very specific and/or highly-tuned
- 3 GPU architecture and execution model are designed for massive parallelism
- 4 GPU efficiency is based on 4 factors:
  - Sufficient parallelism
  - Balanced use of registers and shared memory
  - Efficient DRAM utilization
  - Minimization of divergent branching

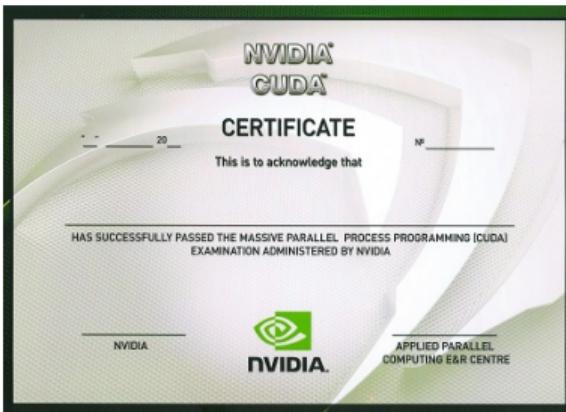
- 1 GPUs offer much higher performance of massively-parallel computations
- 2 You can start using GPUs in your application very easily:
  - Seamless integration into applications with existing parallelism (MPI, OpenMP)
  - Fairly efficient GPU-enabled libraries exist for many types of computational problems
  - Nowadays CUDA is needed only to develop something very specific and/or highly-tuned
- 3 GPU architecture and execution model are designed for massive parallelism
- 4 GPU efficiency is based on 4 factors:
  - Sufficient parallelism
  - Balanced use of registers and shared memory
  - Efficient DRAM utilization
  - Minimization of divergent branching

# GPU / CUDA Training & Certification



## Academic: Parallel & Distributed Computing LAB

- Institute of Computational Science, USI Lugano
- Prof. Dr. Olaf Schenk
- 1 semester (master course), assignments on GPU cluster
- GPU, profiling, optimization, scientific applications case studies
- Contact: [dmitry.mikushin@usi.ch](mailto:dmitry.mikushin@usi.ch)



## Commercial: Advanced GPU computing course

- Applied Parallel Computing LLC & NVIDIA
- 3-5 days, on-site, hands-ons on GPU cluster
- GPU, profiling, optimization, customer-oriented hands-ons
- Register: <http://www.nvidia.de/object/gpu-computing-workshop-registration-form.html>



**Certificate:** granted upon successful pass of comprehensive practical tests